

# Computing Methods for Particle Physics

Bayesian Inference

# Overview

- Today: Bayesian Inference!
- Bayes theorem
- Markov Chain Monte Carlo & Metropolis-Hastings algorithm
- A simple example in Python
- A similar example using PyMC3

# Bayes Theorem


$$P(H|d) = \frac{P(d|H)P(H)}{P(d)}$$

$P(H|d)$ : **posterior**; conditional probability of the hypothesis given the data

$P(d|H)$ : **likelihood**; conditional probability of the data given the hypothesis

$P(H)$ : **prior**; probability of the a particular value for the hypothesis



$P(d)$ : **marginal** distribution; normalisation factor



# Graphically...

$$P(H|d) = \frac{P(d|H)P(H)}{P(d)}$$

Two examples of different priors applied to the same Likelihood



# Likelihood

The **likelihood** introduces our data and our model. It tells us how plausible the data are given the parameters of the model.

**Example:** if we think our data is distributed as a Gaussian and we want to know the mean and standard deviation then we would use a likelihood:

$$p(d|\mu, \sigma) = \frac{1}{\sqrt{2\pi\sigma^2}} \prod_i^N \exp\left(\frac{-(d_i - \mu)^2}{2\sigma^2}\right)$$

Notice that if our data is far away from the distribution suggested by our model then this likelihood is small.

Bayes theorem tells us that this likelihood, along with the **prior**, determines the **posterior** distribution.

# Priors


A **prior** tells us the degree of belief we have in a hypothesis before the data comes in. We must choose the priors for our particular problem carefully.

**Criticism:** A criticism that is often levelled at Bayesian analysis is that the priors are subjective. In general, but for particle physics in particular, the priors aren't all that subjective (e.g. [PDG](#) values are often used).

- Example: a parameter to our model might be the mass of the top particle,  $m_t$ , we might give this a prior of  $\text{norm}(172.9 \text{ GeV}, 0.4 \text{ GeV})$ .

**Note:** When the data is informative the posterior distribution is insensitive to which prior distribution is used. */i.e.* It is only when we lack data that we have to rely on our prior beliefs.

```
%matplotlib inline
import matplotlib.pyplot as plt
import numpy as np
from scipy.stats import norm
import seaborn as sns
sns.set(style='ticks')
top_mass, uncertainty = 172.9, 0.4
x_axis = np.arange(171, 175, 0.001)
y_axis = norm.pdf(x_axis, top_mass, uncertainty)
plt.plot(x_axis, y_axis, color='darkred')
plt.fill_between(x_axis, 0, y_axis,
                 color='darkred', alpha=0.5)
plt.xlabel('top quark mass [GeV]')
plt.ylabel('Density')
plt.show()
```



# Finding the Posterior

In the simple, Gaussian example we could just integrate and solve for the posterior. In general, this is not possible (or we just want a generic solver). Want an approach that can be applied across high dimensional parameter spaces, without taking too long!

In principle we can build an algorithm to just scan the parameter space, computing the likelihood at each point, and multiplying by the prior, to solve for the posterior. But this can be **computationally** prohibitive with large dimensional parameter spaces!

A Markov Chain Monte Carlo (**MCMC**) is an approach that solves this problem by **efficiently** sampling from the posterior distribution.

# Markov Chain Monte Carlo

A **Markov Chain** is the result of any algorithm that takes **only** the current state of the chain,  $X^i$ , step as input to generate a new proposed step,  $X^{i+1}$ . That is

$$p(X^{(i+1)}) = t(X^{(i+1)}|X^{(i)})$$

Where  $t(x|y)$  represents the **transition probability** from one state to another. So the probability that the new state is some particular value depends on the conditional transition probability of the new state given the current state.

We can formulate an MCMC algorithm that has as its equilibrium distribution the posterior from the Bayes formula.



# Metropolis-Hastings

The **Metropolis-Hasting** (MH) algorithm is an MCMC method for sampling the posterior distribution. It builds a **chain** that has visited regions of parameter space with **high posterior probability** relatively more often than those with low posterior probability

## Steps:

1. Start with an initial value (possibly random) of model parameters
2. Choose new proposal parameters (random walk).
3. Accept the proposal with probability  $p_{\text{accept}}$ , based on the ratio of the posteriors.
4. Store into chain if accepted; otherwise store the previous step
  - Note: you always store a value!
5. Repeat

$$p_{\text{accept}} = \min \left[ 1, \frac{\frac{P(x|\mu)P(\mu)}{P(x)}}{\frac{P(x|\mu_0)P(\mu_0)}{P(x)}} = \frac{P(x|\mu)P(\mu)}{P(x|\mu_0)P(\mu_0)} \right]$$

# An MCMC-MH example

Let's look at a simple **example** of the MH algorithm for computing a posterior using an MCMC with as little code as possible (but no MCMC libraries, yet).

It generates a chain for a Gaussian model ( $\mu$  only,  $\sigma$  is fixed) to a data set drawn from a Gaussian to obtain a posterior distribution.

Will point out a few features using this simple [example](#).

# Underflow

$$p(d|\mu, \sigma) = \frac{1}{\sqrt{2\pi\sigma^2}} \prod_i^N \exp\left(\frac{-(d_i - \mu)^2}{2\sigma^2}\right)$$

Beware of machine precision!!

Consider a likelihood that is being computed way out on the tails of the distribution. We might get an extremely small value!

And if we have 1000 data points, using a likelihood of this form will multiple 1000 very small numbers yielding a number too small to be represented as a floating point number. This is known as **underflow**.

The usual method for avoiding this is to take the **negative of the natural log** of the likelihood. So that multiplying many small numbers becomes a sum of reasonable sized negative numbers.

# MH Parameters

## Step size

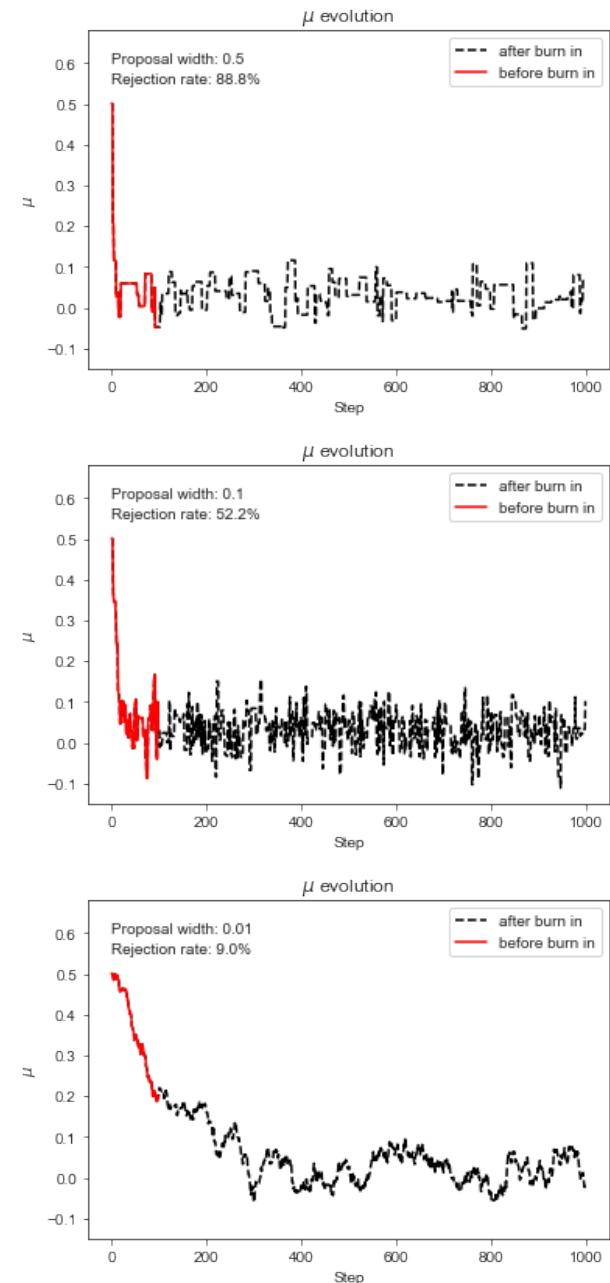
Note that in this simple example we drew the next step from a normal pdf centred around the previous step value and with standard deviation equal to the proposal width

### Autocorrelation!

## Starting value

The starting value takes a certain amount of time to be “forgotten” by the chain. This depends on the step size and the strength of the data.

### Burn-in!

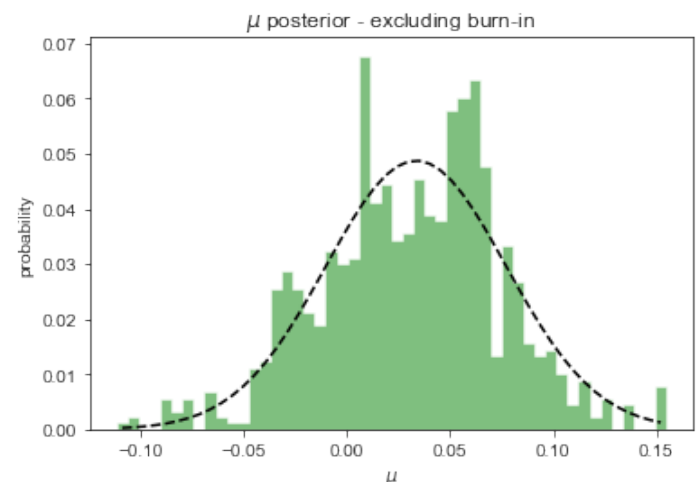
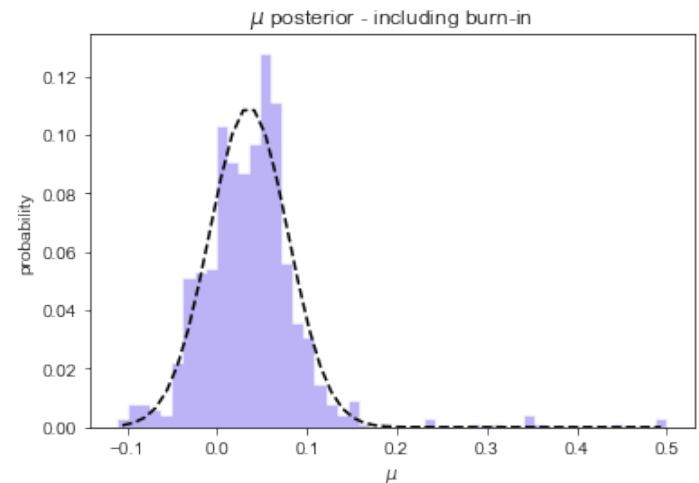
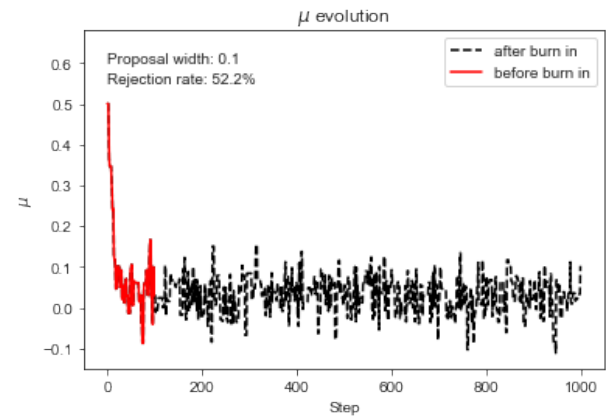


# Burn-in

From the simple example, notice the behaviour of the trace.

The Markov Chain starts off with some memory of the initial state of the chain. This can lead to a bias in our posterior distribution.

We choose some cutoff to throw away earlier portions of the chain. This is known as **burn-in**.





[PyMC3](#) is a python package for Bayesian modelling and machine learning.



- Does the MCMC & MH for you! Along with more advanced implementations!
- Can handle thousands of parameters
- Not slow! Relies on Theano which dynamically compiles code before running MCMC.
- Has a suite of analysis tools and plots to optimise your MCMC.

Let's [revisit](#) our **example**, but now using PyMC3 and also fitting for the standard deviation.

# Uninformative Priors?

Let's revisit the example with lower statistics ( $N=5$ )...

Notice how the shape of the posterior with respect to  $\sigma$  changes drastically:



We used a **uniform** prior on  $\sigma$ . Why is the posterior **biased** upwards in  $\sigma$  when we rely on the prior?

# Uninformative Priors

An **uninformative prior** (or regular, or reference) is one that provides no prior information or bias to the parameter.

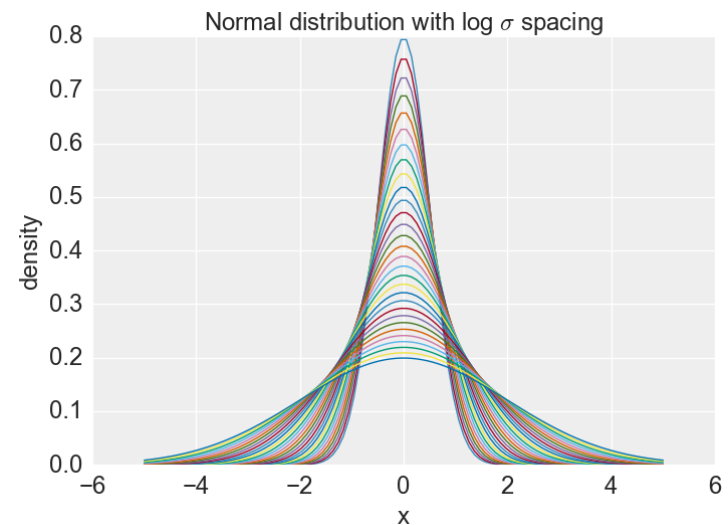

We used a **uniform** prior on sigma, which is not an **uninformative** prior! It biases the likelihood distribution towards higher values of sigma.

A prior that is **invariant under reparameterisation** is known as a **Jeffreys prior**.

For the normal distribution the Jeffreys priors are:

$$\mu \sim 1 \text{ (uniform)}$$

$$\sigma \sim 1/\sigma$$






# Uninformative Priors!

We can use the **Inverse Gamma** distribution to approximate the  $1/x$  behaviour and implement an uninformative prior on  $\sigma$ .

(This also allows us to use built in PyMC3 distributions. You could also code a custom  $1/x$  based distribution, but we will keep things simple.)

**Intuitively** this is reweighting the probability of the  $\sigma$  prior toward lower values.



# Deciding on priors


In practice, you will often be fitting data to models that have well defined priors:

- Particle masses
- Most probable value for Landau distribution ( $dE/dx$  of particle in medium)
- Neutrino oscillation parameters of the PMNS matrix

That is, an uninformative prior is rarely needed: either you have an implicit prior in your problem, or enough data for the prior not to matter. Part of a complete Bayesian inference is **establishing** and **communicating** that one of these is the case.

# Autocorrelation

A chain, particular one generated by the MH algorithm, will tend to have samples correlated with their previous chain entries. An **autocorrelation** plot shows this effect:



This reduces the effective sampling rate of the chain, since we get correlated chain elements. In the past, MCMC practitioners have performed thinning (taking every  $n^{\text{th}}$  chain element) to avoid autocorrelation.

Generally unnecessary. Sufficient to just ensure the chain is long enough.

**Suggestion:** run multiple chains to look for convergence.

# Convergence

We can run multiple chains to test many important features of a chain.

Is our chain long enough?

Is autocorrelation biasing our posterior?

One test, built-in to PyMC3, is the **Gelman-Rubin Test**. This test checks for convergence by comparing the variance of each model parameter between multiple chains. The value of this test, R-hat, should be near 1.

# An alternative approach...

Could also solve this problem with a frequentist approach.

Generally the results will agree if the parameters are all Gaussian. But keep in mind that the two approaches answer different questions.

See [here](#) for a good discussion, using Python examples, of when they differ.

- This really cleared things up for me. Looking at and modifying code has a way of making things more intuitive!

# More advanced inference

We only discussed the basics of MCMC mapping of the posterior distribution using the MH algorithm. There are many more advanced topics you should be aware of and explore depending on your analysis needs:

- Other MCMC samplers
- Parallel chains
- Hierarchical models
- Bayes factors (for comparing one model to another)

# Summary

Bayesian inference is a huge area and we've barely scratched the surface.

Bayes Theorem, Priors, MCMC, MH, PyMC3

You should have enough information to get started on implementing a Bayesian fit to a simple model of your chosen data set.

# Suggested Exercises

1. Read through [Frequentism and Bayesianism: A Python-driven Primer](#) and try to understand the edge cases where these approaches give different results.
2. For the simple MCMC-MH implementation, change the number of data points to 1000 (instead of 100), what happens? Why? How could the code be modified to prevent this?
3. For the PyMC example:
  - a. Explore the strength of the prior relative to the likelihood by varying the size of the generated data set
  - b. Verify that the  $\sigma$  parameter is biased at low statistics. Implement the inverse gamma function as an uninformative prior.